

Language Models

N-Gram

Dr. Uzair Ahmad

Program

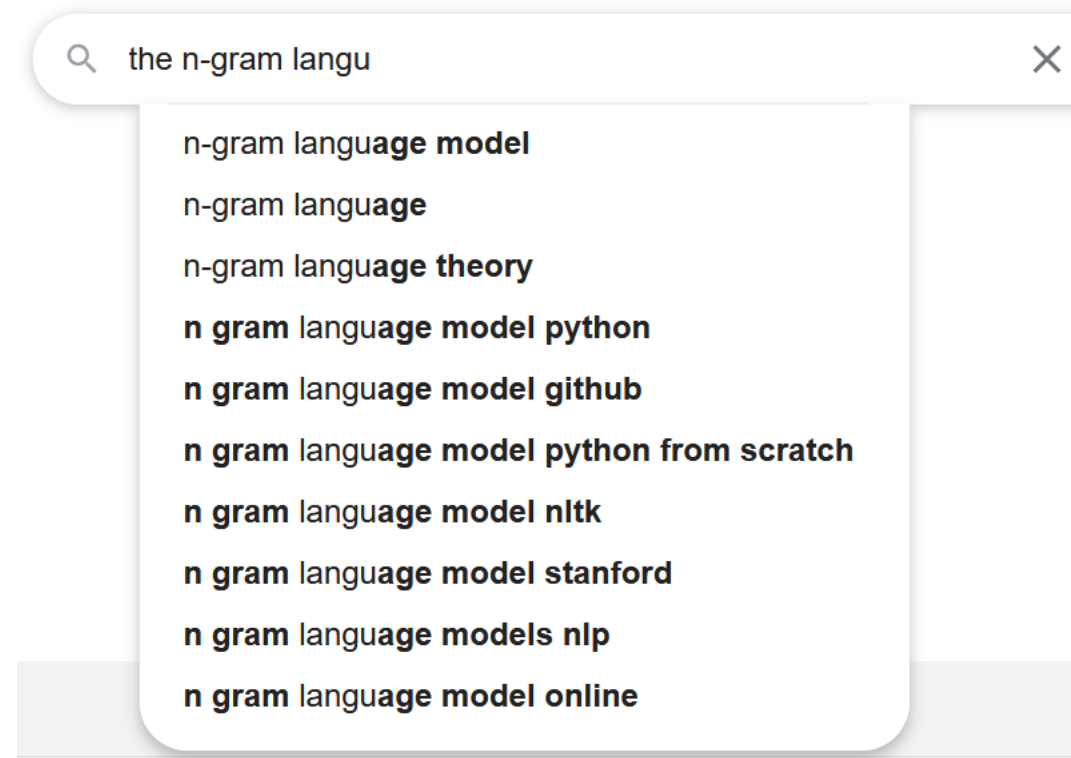
- N-Gram Language Models
 - Predict next word
 - Predict a sentence
- Evaluation of language models

Motivation

- You are uniformly charming!” cried he, with a smile of associating and now and then I bowed and they perceived a chaise and four to wish for.

Random sentence generated from a Jane Austen trigram model

Predicting the next word



Probability of a sequence

- Relative frequency approach

Unbiased estimate $p(\text{Computers are useless, they can only give you answers})$
 $= \frac{\text{count}(\text{Computers are useless, they can only give you answers})}{\text{count}(\text{all sentences ever spoken})}$

Length 9

How many examples are good enough ?

- [Colorless green ideas sleep furiously](https://en.wikipedia.org/wiki/Colorless_green_ideas_sleep_furiously)

https://en.wikipedia.org/wiki/Colorless_green_ideas_sleep_furiously

N-Gram Language Models

- Chain rule refactoring approach

Biased estimate

$$p(\mathbf{w}) = p(w_1, w_2, w_3 \dots w_M) \\ p(w_1)p(w_2|w_1)p(w_3|w_1, w_2) \dots p(w_M|w_{M-1}, \dots, w_1)$$

$$p(\textit{useless} \mid \textit{computers are}) = \frac{\text{count}(\textit{computers are useless})}{\sum_{x \in \mathcal{V}} \text{count}(\textit{computers are } x)} \\ = \frac{\text{count}(\textit{computers are useless})}{\text{count}(\textit{computers are})}.$$

N-Gram Language Models

- The chain rule of probability
 - Product of probabilities of subsequences
 - $p(\mathbf{w}) = p(w_1, w_2, w_3 \dots w_M)$
 - $p(\mathbf{w}) = p(w_1)p(w_2|w_1)p(w_3|w_1, w_2) \dots p(w_M|w_{M-1}, \dots, w_1)$
 - Reverse order
 - $p(\mathbf{w}) = p(w_M)p(w_{M-1}|w_M)p(w_{M-2}|w_{M-1}, w_M) \dots p(w_1|w_2, \dots, w_M)$
 - $\prod_{m=1}^M p(w_m|w_{m-1}, \dots, w_1)$

N-Gram Language Models

$$p(\mathbf{w}) = p(w_1)p(w_2|w_1)p(w_3|w_1, w_2) \dots p(w_M|w_{M-1}, \dots, w_1)$$

$$p(w_M|w_{M-1}, \dots, w_1) \approx p(w_M|w_{M-1}, \dots, w_{M-n+1})$$

$$p(\mathbf{w}) = p(w_1, w_2, w_3 \dots w_M) \approx \prod_{m=1}^M p(w_m|w_{m-1}, \dots, w_{m-n+1})$$

$$\prod_{m=1}^M p(w_m|w_1:w_{m-1}) \approx \prod_{m=1}^M p(w_{m-N+1}|w_{m-1})$$

The Markov assumption

$$\prod_{m=1}^M p(w_m | w_1 : w_{m-1}) \approx \prod_{m=1}^M p(w_{m-N+1} | w_{m-1})$$

$$\begin{aligned} p(I, like, black, coffee) \approx & \\ p(I | \odot) \times & \\ p(like | I) \times & \\ p(black | like) \times & \\ p(coffee | black) \times & \\ p(\otimes | coffee) & \end{aligned}$$

How many **contexts** needed for N-Gram Model ?

N-Gram Language Models

- A Bi-Gram model

$$p(w_m | w_{m-1}) = \frac{\text{count}(w_m, w_{m-1})}{\sum_{w'} \text{count}(w', w_{m-1})} = \frac{\text{count}(w_m, w_{m-1})}{\text{count}(w')}$$

<s> I am Sam </s>

<s> Sam I am </s>

<s> I do not like green eggs and ham </s>

$$P(I | <s>) = \frac{2}{3} = .67 \quad P(\text{Sam} | <s>) = \frac{1}{3} = .33 \quad P(\text{am} | I) = \frac{2}{3} = .67$$

$$P(</s> | \text{Sam}) = \frac{1}{2} = 0.5 \quad P(\text{Sam} | \text{am}) = \frac{1}{2} = .5 \quad P(\text{do} | I) = \frac{1}{3} = .33$$

N-Gram Language Models

- An N-Gram model

$$p(w_m | w_{m-N+1} : w_{m-1}) = \frac{\text{count}(w_{m-N+1} : w_{m-1} w_m)}{\text{count}(w_{m-N+1})}$$

- The hyper-parameter N:
 - **Gorillas** always like to groom **their** friends
 - Too small → High bias:
 - Too big → Large variance

Relative Frequency

N-Gram Language Models

- Large N but ... smoothing for Unknowns

$$p_{smooth}(w_m | w_{m-1}) = \frac{\text{count}(w_m, w_{m-1}) + \alpha}{\sum_{w'} \text{count}(w', w_{m-1}) + |V|\alpha}$$

- Lidstone smoothing
 - Laplace smoothing $\alpha = 1$
 - Jeffreys-Perks law $\alpha = 0.5$

N-Gram Language Models

- Effective counts

$$p_{smooth}(w_m | w_{m-1}) = \frac{count(w_m, w_{m-1}) + \alpha}{\sum_{w'} count(w', w_{m-1}) + |V|\alpha}$$

$$c_i^* = (c_i + \alpha) \frac{M}{M + |V|\alpha}$$

- Discount for each n-gram

$$d_i = \frac{c_i^*}{c_i} = \frac{(c_i + \alpha)}{c_i} \frac{M}{M + |V|\alpha}$$

N-Gram Language Models

- **Smoothing** and **Absolute Discounting**

			Lidstone smoothing, $\alpha = 0.1$		Discounting, $d = 0.1$	
	counts	unsmoothed probability	effective counts	smoothed probability	effective counts	smoothed probability
<i>impropriety</i>	8	0.4	7.826	0.391	7.9	0.395
<i>offense</i>	5	0.25	4.928	0.246	4.9	0.245
<i>damage</i>	4	0.2	3.961	0.198	3.9	0.195
<i>deficiencies</i>	2	0.1	2.029	0.101	1.9	0.095
<i>outbreak</i>	1	0.05	1.063	0.053	0.9	0.045
<i>infirmity</i>	0	0	0.097	0.005	0.25	0.013
<i>cephalopods</i>	0	0	0.097	0.005	0.25	0.013

N-Gram Language Models

- Backoff

$$c^*(i, j) = c(i, j) - d$$
$$P_{\text{Katz}}(i \mid j) = \begin{cases} \frac{c^*(i, j)}{c(j)} & \text{if } c(i, j) > 0 \\ \alpha(j) \times \frac{P_{\text{unigram}}(i)}{\sum_{i': c(i', j)=0} P_{\text{unigram}}(i')} & \text{if } c(i, j) = 0. \end{cases}$$

N-Gram Language Models

- Interpolation

$$\begin{aligned} p_{\text{Interpolation}}(w_m \mid w_{m-1}, w_{m-2}) &= \lambda_3 p_3^*(w_m \mid w_{m-1}, w_{m-2}) \\ &\quad + \lambda_2 p_2^*(w_m \mid w_{m-1}) \\ &\quad + \lambda_1 p_1^*(w_m). \end{aligned}$$

Evaluation of N-Gram Language Models

- Extrinsic
 - Model embedded in Application
- Intrinsic
 - Train-Test Split
 - K-Fold Cross Validation

Evaluation of N-Gram Language Models

- The Perplexity

$$PP(w) = P(w_1, w_2, w_3 \dots w_M)^{-\frac{1}{M}}$$

$$PP(w) = \sqrt[M]{\frac{1}{P(w_1, w_2, w_3 \dots w_M)}}$$

$$PP(w) = \sqrt[M]{\frac{1}{\prod_{i=1}^M P(w_i | w_1, w_2 \dots w_{i-1})}}$$

Summary and Questions

- N-Gram Language Models
 - Predict next word
 - Predict a sentence
- Smoothing relative frequencies
- Evaluation of language models
 - Intrinsic: Perplexity
 - Extrinsic